

WHITE PAPER · AI-POWERED MANAGEMENT ANALYTICS

TurboHUB

Shortening the “First Mile” of Management Reporting

Reliable management answers from fragmented enterprise data in minutes with TurboHUB

Executive Summary

Management decisions can only move as fast as the data behind them. In many organizations, accounting, ERP, payroll, CRM, warehouse systems and Excel models operate in isolation. A single executive question turns into days of data extraction, comparison and reconciliation by analysts.

TurboHUB is a continuously refreshed operational data layer built on top of existing systems. Data from different sources is unified by AI into one consistent layer. Users ask questions in natural language, and the system returns a structured, actionable answer within minutes.

This paper outlines the root cause of the problem, how TurboHUB works, its impact on each department, the available deployment options and the infrastructure requirements for decision makers and IT teams.

CONTENTS

- 01 Executive Summary
- 02 The Problem: Fragmented Data and Hidden Reporting Costs
- 03 The Solution: How TurboHUB Works
- 04 Impact by Department and Business Benefits
- 05 Deployment Models and Data Security
- 06 Hardware and Infrastructure Requirements
- 07 Implementation Roadmap and Contact

Highlights

5+ days → 2 min

Request-to-report cycle time

100%

Elimination of report customization costs

2+ FTE

Analyst capacity freed up

Figures reflect target values stated in TurboHUB product material. Actual results depend on data sources, data quality and project scope.

Who is TurboHUB for?

- ✓ Companies consolidating data from multiple systems for management reporting
- ✓ Teams where finance and analysts spend most of their time preparing data
- ✓ Organizations where report requests pile up in the IT backlog
- ✓ Organizations that want to use AI without moving data outside their premises

01

The Problem: Fragmented Data and Hidden Reporting Costs

Companies using BI tools or Excel for management reporting find that the problem does not start at analysis but at data preparation, the “first mile” of analytics.

5–7**data sources**

Manual reconciliation across accounting, payroll, CRM, warehouse systems, external services and Excel models

5–15%**discrepancy**

Variation of the same metric from one system to another

3–5 days**every month**

Business days spent on reconciliation instead of interpretation

Three types of inefficient costs

01

IT Department

Instead of its core mandate of infrastructure development, IT is tied up customizing and maintaining reports. Every new report dimension means a ticket, a 2–3 week wait and a growing backlog.

02

Analysts and Finance Managers

Spend 70–80% of their time not on analysis, but on manually extracting, comparing and reconciling data from multiple sources.

03

Accounting Department


Adds extra analytical dimensions to primary documents purely for management reporting, not for statutory accounting, which increases workload and error risk.

The result: the organization's most valuable analytical capacity is spent preparing data instead of interpreting it, and decisions are delayed.

02

The Solution: How TurboHUB Works

TurboHUB is a continuously updated data layer built on top of your existing systems, without replacing them, combined with an AI interface that lets you query that layer in natural language.



How it works in four steps

01

Ask

The user asks a question in natural language; no special skills or training are required.

02

Fresh data

TurboHUB keeps an operational data layer fed by enterprise systems continuously up to date.

03

Unify

Data from different sources is matched and unified by AI into a single layer.

04

Answer

The system returns a structured, comparable and actionable output.

Example executive questions

"Show revenue breakdown by counterparty for last month."

"Which cost items deviate most from budget this quarter?"

Questions are illustrative; available dimensions depend on the connected data sources.

Analytics via indirect indicators

TurboHUB can extract analytical insight not only from structured fields but also from text fields and related documents. Accounting teams no longer need to add tags purely for reporting purposes.

03

Impact by Department and Business Benefits

TurboHUB introduces a different operating model at the three points where inefficient costs arise.

DEPARTMENT	TODAY	WITH TURBOHUB
IT Department	⊗ An IT ticket for every new report dimension ⊗ 2–3 weeks lead time for customization ⊗ Backlog of pending business requests	✓ New report dimension in 2–4 minutes without IT ✓ One consistent logic across all reports
Analysts and Finance	⊗ 3–5 days of manual consolidation each month ⊗ Prolonged reconciliation of discrepancies ⊗ 70–80% of time on routine tasks	✓ Natural language question, answer in minutes ✓ Frees up 1.5–2 FTE of analyst capacity
Accounting	⊗ Manual analytical tagging of primary documents ⊗ Increased data entry workload ⊗ High risk of human error	✓ AI analytics from text fields and documents ✓ Team focuses on statutory compliance

Key business benefits

5+ days → 2 min

Shorter request-to-report cycle

100%

Elimination of report customization costs

2+ FTE

Analyst capacity freed up

Less key-person dependency

Reporting knowledge no longer sits with individual employees; business logic lives in the system.

Faster decisions

Where a 24-hour delay means lost profit, answers arrive in minutes.

Ease of use

No specialized technical skills or training are required to use the system.

Figures reflect target values stated in TurboHUB product material. Actual results depend on data sources, data quality and project scope.

04

Deployment Models and Data Security

TurboHUB can be deployed in three models depending on your security policies and infrastructure.

01

On-premise

Installed on the organization's own physical servers or virtual machines.

- ✓ Data stays within your infrastructure
- ✓ Can run with local language models (LLMs)
- ✓ Container-based; virtualization support required

02

Hardware appliance

If the required servers are not available, TurboHUB is delivered pre-installed and configured on dedicated hardware.

- ✓ Compact AI system based on NVIDIA DGX Spark
- ✓ Placed inside your infrastructure
- ✓ Fast time to value

03

Cloud (SaaS)

If new systems cannot be hosted internally, TurboHUB is provided as a service on dedicated servers.

- ✓ Local LLMs on GPU servers
- ✓ or CPU servers with cloud LLM subscriptions
- ✓ Configuration sized to requirements

Data security principles



Data sovereignty

In on-premise and appliance models, data never leaves the organization's own infrastructure.



Local AI models

Language models can run in-house, so sensitive data does not have to be sent to third-party cloud services.



No changes to source systems

TurboHUB sits on top of source systems and does not alter accounting or ERP processes.



Controls defined in scoping

Access rights, connected sources and data scope are defined with your IT and security teams at project start.

Hardware and Infrastructure Requirements

The configurations below are indicative. Final sizing is completed during scoping based on user count, data volume and the selected language model.

On-premise deployment profiles

PROFILE	USERS	CPU	MEMORY	STORAGE	GPU OPTIONS
Minimum	1–2	Intel Core i7 (12th gen or newer)	128 GB	4 TB SSD	NVIDIA RTX 3060 / 4090 / 5090
Pilot / test	1–10	Intel Core i9 (14th gen)	128 GB	8 TB SSD	NVIDIA Tesla P40 24 GB / A40 48 GB / A100 80 GB
Enterprise	100+	Intel Xeon Platinum	512 GB	16 TB SSD	NVIDIA A100 80 GB / H100 80 GB / H200 141 GB / 4× L40 48 GB

PCIe 5.0 and DDR5 support are recommended for all profiles. The physical server or virtual machine must support virtualization for container-based deployment.

Reference server architectures

Mid-scale	2U server · 2× Intel Xeon (18 cores) · 512 GB DDR5 · 4× 960 GB SSD (RAID) · 2× NVIDIA L40S 48 GB · 10/25 GbE networking · redundant power
High-scale	6U server · 2× Intel Xeon Platinum (64 cores) · 2 TB DDR5 · 5× 3.84 TB NVMe (RAID) · NVIDIA HGX H200 8-GPU (141 GB) · 10/25 GbE networking · redundant power

Hardware appliance

Single system	1× NVIDIA DGX Spark compact AI supercomputer
Two-node cluster	2× NVIDIA DGX Spark linked via 200G QSFP56; models up to 405 billion parameters

Cloud (SaaS) server profiles

GPU · long-context production	NVIDIA RTX PRO 6000 96 GB · AMD Ryzen 9 9950X (16 cores) · 192 GB DDR5 · 2× 2 TB NVMe
GPU · high-load inference	4× NVIDIA A100 40 GB · 2× Intel Xeon Gold (48 cores) · 256 GB ECC · NVMe + SSD
CPU · standard	AMD Ryzen 9 9950X (16 cores) · 192 GB DDR5 ECC · 1 TB + 2× 1.92 TB NVMe · with cloud LLM
CPU · high-performance	2× Intel Xeon 6710E (128 cores) · 256 GB DDR5 ECC · 2× 1.92 TB NVMe · with cloud LLM




Pricing is prepared as a project-specific proposal based on license scope, deployment model and hardware configuration.

06

Implementation Roadmap

TurboHUB is rolled out in low-risk, measurable steps.



About AI Kitchen

AI Kitchen helps organizations automate repeatable business processes with RPA, document automation, data processing and AI-supported workflows. With TurboHUB, it extends this approach to management reporting and analytics, giving decision makers reliable data within minutes.

Let's start with a demo

In 15 minutes, we'll show you how TurboHUB works and the impact it can create in your organization.

ET

Elena Türkmen

TurboHUB Contact · AI Kitchen

elena.turkmen@aikitchen.com.tr
aikitchen.com.tr